

A-star algorithm

[A* Search Algorithm - GeeksforGeeks](#)

[A* search algorithm - Wikipedia](#)



一个带有启发式指引 (greedy) 的两点间最短路的搜索的Dijkstra算法。

背景

- 一种图遍历的路径搜索算法;
- 给定一个带权图, 不是处理单个源节点到所有节点的最短路问题, 是处理单个源节点到某一个确定的目标节点的**最短路**;

Compared to Dijkstra's algorithm, the A* algorithm only finds the shortest path **from a specified source to a specified goal**, and not the shortest-path tree from a specified source to all possible goals.

- 时间复杂度 $O(|E|\log|V|)$, 空间复杂度 $O(|V|)$;
- 可以视为Dijkstra算法的拓展; 其使用了一个启发式函数指导其搜索;

This is a necessary trade-off for using a specific-goal-directed heuristic.

- 虽然每次只选择当前的最小成本路径, 不过该greedy策略最终可以得到一个全局最优的结果;

算法内容

<https://en.wikipedia.org/wiki/File:Astarpathfinding.gif>

- It does this by maintaining a **tree** of paths originating at the start node and extending those paths one edge at a time until the goal node is reached.
- 每次主循环, 算法需要挑选一条最小成本的路径进行extend; 代价函数分为两部分, 路径的cost和把路径延伸到目标的估计cost。

$$f(n) = g(n) + h(n)$$

- n 是路径上的下一个节点, $g(n)$ 是从起始节点到 n 的路径成本, $h(n)$ 是一个启发式函数, 用于估计从 n 到目标的最小成本路径的成本 (the cost of the cheapest path from n to the goal) ;

We define 'g' and 'h' as simply as possible below

g = the **movement cost** to move from the starting point to a given square on the grid, following the path generated to get there.

h = the estimated movement cost to move from that given square on the grid to the final destination. This is often referred to as the heuristic, which is nothing but a kind of **smart guess**.

实现

- A* 的典型实现使用优先队列 (作为frontier) 来执行最小 (估计) 成本节点的选择以进行扩展。
- 在算法的每一步, 具有最低 $f(x)$ 值的节点从队列中pop, 其邻居的 f 和 g 值相应更新, 并将这些邻居添加到队列中。
- 终止条件: pop的节点恰好为目标节点。
- 被访问过的节点, 都存储了对父节点的引用, 因此从目标节点倒推回去就能得到最短路;
- $h(x)$ 可以被设置为是物理距离, 比如L1距离, L2距离;

优先队列的实现通常是斐波那契堆, 复杂度为 $O(\log V)$;

伪代码

```
function reconstruct_path(cameFrom, current)
    total_path := {current}
    while current in cameFrom.Keys:
        current := cameFrom[current]
        total_path.prepend(current)
    return total_path

// A* finds a path from start to goal.
// h is the heuristic function. h(n) estimates the cost to reach goal from node n.
function A_Star(start, goal, h)
    // The set of discovered nodes that may need to be (re-)expanded.
    // Initially, only the start node is known.
    // This is usually implemented as a min-heap or priority queue rather than a hash-set.
    openSet := {start}

    // For node n, cameFrom[n] is the node immediately preceding it on the cheapest path from the start
    // to n currently known.
    cameFrom := an empty map

    // For node n, gScore[n] is the cost of the cheapest path from start to n currently known.
    gScore := map with default value of Infinity
    gScore[start] := 0

    // For node n, fScore[n] := gScore[n] + h(n). fScore[n] represents our current best guess as to
    // how cheap a path could be from start to finish if it goes through n.
    fScore := map with default value of Infinity
    fScore[start] := h(start)

    while openSet is not empty
        // This operation can occur in O(Log(N)) time if openSet is a min-heap or a priority queue
        current := the node in openSet having the lowest fScore[] value
        if current = goal
            return reconstruct_path(cameFrom, current)

        openSet.Remove(current)
        for each neighbor of current
            // d(current,neighbor) is the weight of the edge from current to neighbor
            // tentative_gScore is the distance from start to the neighbor through current
            tentative_gScore := gScore[current] + d(current, neighbor)
            if tentative_gScore < gScore[neighbor]
                // This path to neighbor is better than any previous one. Record it!
                cameFrom[neighbor] := current
                gScore[neighbor] := tentative_gScore
                fScore[neighbor] := tentative_gScore + h(neighbor)
                if neighbor not in openSet
                    openSet.add(neighbor)

    // Open set is empty but goal was never reached
    return failure
```

一个很清晰的演示: <https://upload.wikimedia.org/wikipedia/commons/9/98/AstarExampleEn.gif>

难点

- 如何获取有效的启发式函数? path的cost是非常直接就能得到, 但是到目标的距离是一个并不显然的估计值;
- 一个错误的、过于简单的启发式函数, 可能会造成误导的效果, 比如在目标点之前有一堵墙, 那一开始直接朝着目标点探索并不是最优的路径;
- 核心: 对于目标距离的估计 (或者是对未来reward的期望, 对于正确得到答案的概率, anyway随便怎么样都是同一套) 如何基于当前的state进行估计? 回归到经典的RL的问题, 对于state-value/action-value的估计;

- 对于抽象概念的距离，譬如“语义”，“数学推理的正确性”，这种距离如何设计？如何朝着有效的方向进发？如何设计模型进行近似计算？